

Exploring Pandas: Common Data Operations

Welcome to this Jupyter Notebook! 🚀 In this notebook, you'll practice some of the most commonly used operations in the **Pandas** library using **two datasets**:

1. `../data/students.csv` (CSV)
2. `../data/enrollments.json` (JSON)

These files should be placed in the same folder as this notebook. By the end, you'll have a strong grasp of common data manipulation tasks, and you'll even merge these two datasets on a common key.

Before starting, make sure you have **Pandas** installed. (It should come preinstalled in Anaconda!)

If pandas is not installed, follow the instructions below.

Checking if Pandas is Installed in Your Conda Environment

Before proceeding, check if Pandas is installed in your Conda environment by running the following command in a **Jupyter Notebook** cell:

```
In [33]: import pandas as pd  
print(pd.__version__)
```

2.0.3

If this runs without errors and prints a version number, Pandas is installed. If you see an **ImportError**, install Pandas using one of the following methods:

For Conda Users (Recommended)

Run this in your terminal or Anaconda Prompt:

```
conda install pandas
```

Using Conda-Forge (If Needed)

If you encounter issues, you can install Pandas from **Conda-Forge**, a community-maintained repository with up-to-date packages:

```
conda install -c conda-forge pandas
```

For Pip Users

If you're using a virtual environment outside Conda, install Pandas via Pip:

```
pip install pandas
```


Now, let's dive in!

1. Load a CSV file into a Pandas DataFrame

First, let's **import Pandas** and load the datasets. Two datasets have been prepared for you:

- `students.csv`
- `enrollments.json`

You will use these two datasets for the following challenges.

 **Hint:** If the file is in the same directory as your notebook, you can just use the filename. Otherwise, provide the full file path.

```
In [41]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [43]: # Load students.csv
students_df = pd.read_csv("../data/students.csv")

# Load enrollments.json
enrollments_df = pd.read_json("../data/enrollments.json")

# Display the first few rows of both datasets
print(students_df.head())
print(enrollments_df.head())
```

	student_id	First_Name	last_name	Birthdate	gender	majorField	\
0	STU001	John	Doe	4/12/1998	M	Computer Science	
1	STU002	Maria	Gonzalez	9/5/1997	F	Biology	
2	STU003	Priya	Patel	1/23/1999	F	Engineering	
3	STU004	Alex	Johnson	12/15/1996	M	Mathematics	
4	STU005	Emily	Smith	7/30/2000	F	Physics	

	admission_year	current	gpa	contact_email	mobile	number	\
0	2020		3.5	john.doe@example.com	-655.0		
1	2019		3.8	maria.gonzalez@example.com	-656.0		
2	2021		3.7	priya.patel@example.com	-657.0		
3	2018		3.2	alex.johnson@example.com	NaN		
4	2022		3.9	emily.smith@example.com	-659.0		

	home_city	HOME	COUNTRY
0	Tampa		USA
1	London		UK
2	Mumbai		India
3	Sydney		Australia
4	New York		USA

	enrollment_id	stud_ref_id	subject_code	course_title	\
0	ENR-STU001-MATH101-AB12	STU001	MATH101	Calculus I	
1	ENR-STU002-ENG201-CD34	STU002	ENG201	English Literature	
2	ENR-STU003-CS301-EF56	STU003	CS301	Programming 101	
3	ENR-STU004-HIST105-GH78	STU004	HIST105	World History	
4	ENR-STU005-CHEM110-IJ90	STU005	CHEM110	Organic Chemistry	

	instructor_name	enroll_count	term_offered	course_fee	final_result	\
0	Dr. Smith	25	Spring 2026	1000	A	
1	Prof. Brown	26	Fall 2025	1100	B	
2	Dr. Taylor	27	Spring 2026	1200	C+	
3	Dr. Lee	28	Fall 2025	1300	B+	
4	Dr. Kumar	29	Spring 2026	1400	B-	


	attend_percentage	date_enrolled
0	71	2026/01/15
1	72	2025-09-01
2	73	2026/01/15
3	74	2025/09/01
4	75	2026-01-15

In []:

In [3]: # your code here

2. View the First and Last Few Rows of Each DataFrame

Check out how your data looks. One method previews the first few records, while another method previews the last few. You can specify the number of rows you want to see by explicitly passing an integer argument.

 **Tip:** This is a great time to confirm that columns loaded correctly and to spot any obvious data issues (strange values, mismatched columns, etc.).

```
In [45]: print(students_df.head())
print(enrollments_df.head())
```

```

student_id First_Name last_name Birthdate gender majorField \
0   STU001      John      Doe   4/12/1998      M  Computer Science
1   STU002    Maria  Gonzalez   9/5/1997      F      Biology
2   STU003    Priya    Patel   1/23/1999      F    Engineering
3   STU004     Alex   Johnson  12/15/1996      M    Mathematics
4   STU005     Emily     Smith   7/30/2000      F      Physics

admission_year  current  gpa      contact_email  mobile number \
0           2020        3.5  john.doe@example.com    -655.0
1           2019        3.8  maria.gonzalez@example.com -656.0
2           2021        3.7  priya.patel@example.com    -657.0
3           2018        3.2  alex.johnson@example.com    NaN
4           2022        3.9  emily.smith@example.com    -659.0

home_city HOME COUNTRY
0   Tampa      USA
1  London      UK
2  Mumbai    India
3  Sydney  Australia
4  New York      USA

enrollment_id stud_ref_id subject_code      course_title \
0  ENR-STU001-MATH101-AB12    STU001    MATH101      Calculus I
1  ENR-STU002-ENG201-CD34    STU002    ENG201  English Literature
2  ENR-STU003-CS301-EF56    STU003    CS301    Programming 101
3  ENR-STU004-HIST105-GH78    STU004    HIST105    World History
4  ENR-STU005-CHEM110-IJ90    STU005    CHEM110    Organic Chemistry

instructor_name  enroll_count  term_offered  course_fee  final_result \
0   Dr. Smith          25  Spring 2026      1000          A
1  Prof. Brown          26   Fall 2025      1100          B
2   Dr. Taylor          27  Spring 2026      1200         C+
3    Dr. Lee           28   Fall 2025      1300         B+
4   Dr. Kumar          29  Spring 2026      1400         B-

attend_percentage  date_enrolled
0                71   2026/01/15
1                72   2025-09-01
2                73   2026/01/15
3                74   2025/09/01
4                75   2026-01-15
```

```
In [47]: print(students_df.tail())
print(enrollments_df.tail())
```

	student_id	First_Name	last_name	Birthddate	gender	majorField	\
98	STU096	Victoria	Ortiz	8/2/1996	F	Mathematics	
99	STU097	Julian	Foster	9/3/1998	M	Physics	
100	STU098	Lucy	Ramirez	10/4/1997	F	Chemistry	
101	STU099	Isaiah	Kim	11/5/1999	M	Economics	
102	STU100	Amelia	Lopez	12/6/1996	F	History	

	admission_year	current	gpa	contact_email	mobile	number	\
98	2020		3.7	victoria.ortiz@example.com		-750.0	
99	2022		3.8	julian.foster@example.com		-751.0	
100	2019		3.9	lucy.ramirez@example.com		-752.0	
101	2021		3.2	isaiah.kim@example.com		-753.0	
102	2020		3.3	amelia.lopez@example.com		-754.0	

	home_city	HOME	COUNTRY
98	Concord		USA
99	Toledo		USA
100	St. Louis		USA
101	Orlando		USA
102	Raleigh		USA

	enrollment_id	stud_ref_id	subject_code	course_title
96	ENR-STU097-MATH101-MN34	STU097	MATH101	Calculus
97	ENR-STU098-ENG201-OP56	STU098	ENG201	English Literature
98	ENR-STU099-CS301-QR78	STU099	CS301	Programming 101
99	ENR-STU100-HIST105-ST90	STU100	HIST105	World History
100	ENR-STU004-HIST999-DUP1	STU004	HIST999	Ancient Civilizations


	instructor_name	enroll_count	term_offered	course_fee	final_result	\
96	Dr. Smith	30	Spring 2026	1500	A	
97	Prof. Brown	31	Fall 2025	1600	B	
98	Dr. Taylor	25	Spring 2026	1000	C+	
99	Dr. Lee	26	Fall 2025	1100	B+	
100	Prof. Brown	99	Fall 2023	NaN		

	attend_percentage	date_enrolled
96	83	2026-01-15
97	84	2025-09-01
98	85	2026-01-15
99	86	2025-09-01
100	67	2023-09-01

```
In [7]: # your code here
```

3. Check the Shape of Each DataFrame

To understand the **size** of your dataset(s), use the attribute that returns (number_of_rows, number_of_columns) .

 **Tip:** Note any big differences in row counts that might affect merging later.


```
In [49]: print("Students DataFrame Shape:", students_df.shape)
         print("Enrollments DataFrame Shape:", enrollments_df.shape)
```

```
Students DataFrame Shape: (103, 12)
Enrollments DataFrame Shape: (101, 11)
```

4. Get a Summary of Each DataFrame

Explore one or two approaches that provide:

- **Column names**
- **Data types**
- **Basic statistics about numerical columns**
- **Number of non-null values**

 **Tip:** One approach might give an overview of columns and data types; another might summarize numerical columns. This step helps you detect columns that might need cleaning.

```
In [51]: print(students_df.info())
         print(enrollments_df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 103 entries, 0 to 102
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   student_id            103 non-null    object
1   First_Name            103 non-null    object
2   last_name             103 non-null    object
3   Birthddate            103 non-null    object
4   gender                103 non-null    object
5   majorField            102 non-null    object
6   admission_year        103 non-null    object
7   current gpa           95 non-null     float64
8   contact_email         103 non-null    object
9   mobile_number         97 non-null     float64
10  home_city             103 non-null    object
11  HOME_COUNTRY          103 non-null    object
dtypes: float64(2), object(10)
memory usage: 9.8+ KB
None

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 101 entries, 0 to 100
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   enrollment_id         101 non-null    object
1   stud_ref_id           101 non-null    object
2   subject_code          101 non-null    object
3   course_title          101 non-null    object
4   instructor_name       101 non-null    object
5   enroll_count          101 non-null    int64
6   term_offered          101 non-null    object
7   course_fee            100 non-null    object
8   final_result          101 non-null    object
9   attend_percentage     101 non-null    int64
10  date_enrolled         101 non-null    object
dtypes: int64(2), object(9)
memory usage: 8.8+ KB
None
```

```
In [53]: print(students_df.describe())
         print(enrollments_df.describe())
```


	current_gpa	mobile_number
count	95.000000	97.000000
mean	3.560000	-688.381443
std	0.230817	150.548334
min	3.100000	-754.000000
25%	3.400000	-727.000000
50%	3.600000	-703.000000
75%	3.750000	-678.000000
max	3.900000	753.000000

	enroll_count	attend_percentage
count	101.000000	101.000000
mean	28.653465	79.633663
std	7.350423	6.009530
min	25.000000	67.000000
25%	26.000000	75.000000
50%	28.000000	80.000000
75%	30.000000	85.000000
max	99.000000	90.000000

In [13]: *# your code here*

5. Check for Missing Values

Determine if your dataset has any missing or null values by **counting** them. Notice which columns have many missing entries and plan how to handle them.

 **Tip:** Some columns might look present but contain empty strings. Identify them if possible.

```
In [55]: print(students_df.isnull().sum())
print(enrollments_df.isnull().sum())
```



```

student_id      0
First_Name      0
last_name       0
Birthddate      0
gender          0
majorField      1
admission_year  0
current_gpa     8
contact_email   0
mobile_number   6
home_city       0
HOME_COUNTRY    0
dtype: int64
enrollment_id   0
stud_ref_id     0
subject_code    0
course_title    0
instructor_name 0
enroll_count    0
term_offered    0
course_fee      1
final_result    0
attend_percentage 0
date_enrolled   0
dtype: int64

```

In [15]: `# your code here`

In [16]: `# your code here`

6. Rename Columns for Clarity and Consistency

Some columns may have **spaces** or **capitalization** that complicates your analysis. For example, if you see "current_gpa" or "First_Name", consider renaming them (e.g., "current_gpa", "first_name") for ease of use.

 **Tip:** Consistent naming conventions help minimize typos and KeyErrors.

```

In [57]: students_df.rename(columns={
        "First_Name": "first_name",
        "last_name": "last_name",
        "Birthddate": "birthdate",
        "current_gpa": "current_gpa"
    }, inplace=True)

print(students_df.head())

```

	student_id	first_name	last_name	birthdate	gender	majorField	\
0	STU001	John	Doe	4/12/1998	M	Computer Science	
1	STU002	Maria	Gonzalez	9/5/1997	F	Biology	
2	STU003	Priya	Patel	1/23/1999	F	Engineering	
3	STU004	Alex	Johnson	12/15/1996	M	Mathematics	
4	STU005	Emily	Smith	7/30/2000	F	Physics	

	admission_year	current_gpa	contact_email	mobile	number	\
0	2020	3.5	john.doe@example.com		-655.0	
1	2019	3.8	maria.gonzalez@example.com		-656.0	
2	2021	3.7	priya.patel@example.com		-657.0	
3	2018	3.2	alex.johnson@example.com		NaN	
4	2022	3.9	emily.smith@example.com		-659.0	

	home_city	HOME	COUNTRY
0	Tampa		USA
1	London		UK
2	Mumbai		India
3	Sydney		Australia
4	New York		USA


In [18]: *# your code here*


In [19]: *# your code here*

In [20]: *# your code here*

7. Convert Data Types Where Needed

Check which columns should be numeric or datetime. Columns like `admission_year` or `course_fee` might be read as **strings** by default. Convert them to numerical or date formats if necessary.

 **Tip:** Make sure you handle errors gracefully (e.g., set `errors='coerce'` to turn invalid entries into NaN).

In [103]: *# Convert admission_year to numeric*
`students_df["admission_year"] = pd.to_numeric(students_df["admission_year"],`
`print(students_df.dtypes) #  Output the data types after conversion`

```
student_id      object
first_name      object
last_name       object
birthdate       datetime64[ns]
gender          object
majorField      object
admission_year  float64
current_gpa     float64
contact_email   object
home_city       object
HOME_COUNTRY    object
full_name       object
dtype: object
```

```
In [105... # Convert birthdate to datetime
students_df["birthdate"] = pd.to_datetime(students_df["birthdate"], errors='
print(students_df.head()) # ✅ Output first few rows to verify changes
```

	student_id	first_name	last_name	birthdate	gender	majorField	\
0	STU001	John	Doe	1998-04-12	M	Computer Science	
1	STU002	Maria	Gonzalez	1997-09-05	F	Biology	
2	STU003	Priya	Patel	1999-01-23	F	Engineering	
3	STU004	Alex	Johnson	1996-12-15	M	Mathematics	
4	STU005	Emily	Smith	2000-07-30	F	Physics	

	admission_year	current_gpa	contact_email	home_city	\
0	2020.0	3.5	john.doe@example.com	Tampa	
1	2019.0	3.8	maria.gonzalez@example.com	London	
2	2021.0	3.7	priya.patel@example.com	Mumbai	
3	2018.0	3.2	alex.johnson@example.com	Sydney	
4	2022.0	3.9	emily.smith@example.com	New York	

	HOME COUNTRY	full_name
0	USA	John Doe
1	UK	Maria Gonzalez
2	India	Priya Patel
3	Australia	Alex Johnson
4	USA	Emily Smith

8. Fill Missing Values with a Specified Value or Method

Instead of **dropping** missing values, consider **replacing** them. For instance:

- A string like "Unknown" for missing text
- A **mean** or **median** for missing numeric columns
- A **forward** or **backward fill** if appropriate

```
In [129... print(students_df.columns) # ✅ Verify column names
```


```
Index(['student_id', 'first_name', 'last_name', 'birthdate', 'gender',
      'majorField', 'admission_year', 'current_gpa', 'contact_email',
      'home_city', 'HOME COUNTRY', 'full_name'],
      dtype='object')
```


```
In [135... print(students_df.isnull().sum()) # ✅ Ensure missing values are handled
```

```
student_id      0
first_name      0
last_name       0
birthdate       1
gender          0
majorField      0
admission_year  1
current_gpa     0
contact_email   0
home_city       0
HOME_COUNTRY    0
full_name       0
dtype: int64
```

9. Drop Rows or Columns with Missing Values (If Needed)

After considering which values can be filled, you might choose to **remove** rows or columns that are missing too much data or can't be fixed.

 **Tip:** Decide carefully and confirm you don't need the dropped information. Use `inplace=True` or keep a separate DataFrame if you want to preserve the original data.

```
In [137... print(students_df.columns) #  Verify column names

Index(['student_id', 'first_name', 'last_name', 'birthdate', 'gender',
      'majorField', 'admission_year', 'current_gpa', 'contact_email',
      'home_city', 'HOME_COUNTRY', 'full_name'],
      dtype='object')
```

```
In [30]: # your code here
```

```
In [31]: # your code here
```

10. Filter Rows Based on a Condition

Now that columns like `admission_year` and `course_fee` (or `current_gpa`) are numeric, experiment with filtering. For example:

- Students whose `admission_year` is after a certain date
- Enrollments for Spring 2026

```
In [83]: recent_students = students_df[students_df["admission_year"] > 2020]
print(recent_students.head())
```

	student_id	first_name	last_name	birthdate	gender	majorField	\
2	STU003	Priya	Patel	1999-01-23	F	Engineering	
4	STU005	Emily	Smith	2000-07-30	F	Physics	
5	STU006	Robert	Brown	1995-03-20	M	Unknown	
8	STU009	Sophia	Lopez	1999-09-30	F	Computer Science	
12	STU013	Michael	Miller	1999-11-02	M	Physics	

	admission_year	current_gpa	contact_email	home_city	\
2	2021.0	3.7	priya.patel@example.com	Mumbai	
4	2022.0	3.9	emily.smith@example.com	New York	
5	2023.0	3.3	robert.brown@example.com	Los Angeles	
8	2021.0	3.7	sophia.lopez@example.com	Dublin	
12	2021.0	3.7	michael.miller@example.com	Toronto	

	HOME COUNTRY
2	India
4	USA
5	USA
8	Ireland
12	Canada

In [33]: `# your code here`

In [34]: `# your code here`

In [35]: `# your code here`

11. Select Specific Columns from Each DataFrame

Often, you don't need all columns at once. For instance, you might extract only:

- "student_id", "First_Name", "last_name", and "current_gpa" from students.csv
- "stud_ref_id", "course_title", "instructor_name", "course_fee" from enrollments.json

In [85]: `students_subset = students_df[["student_id", "first_name", "last_name", "current_gpa"]]`
`print(students_subset.head())`

	student_id	first_name	last_name	current_gpa
0	STU001	John	Doe	3.5
1	STU002	Maria	Gonzalez	3.8
2	STU003	Priya	Patel	3.7
3	STU004	Alex	Johnson	3.2
4	STU005	Emily	Smith	3.9

In [37]: `# your code here`

In [38]: `# your code here`

In [39]: `# your code here`

12. Sort the DataFrame by One or More Columns

Sorting can help you identify which records have the highest or lowest values. For example:

- Sort the **students** DataFrame by "current_gpa" in descending order
- Sort the **enrollments** DataFrame by "course_fee" in ascending order

 **Tip:** You can sort by multiple columns if needed.

```
In [87]: sorted_students = students_df.sort_values(by="current_gpa", ascending=False)
print(sorted_students.head())
```

	student_id	first_name	last_name	birthdate	gender	majorField	\
31	STU032	Grace	Moore	1996-02-28	F	Mathematics	
68	STU066	Victoria	Morris	1997-02-02	F	Chemistry	
84	STU082	Aurora	Rivera	1997-06-18	F	Chemistry	
92	STU090	Molly	Barnes	1997-02-26	F	Chemistry	
76	STU074	Natalie	Ross	1997-10-10	F	Chemistry	

	admission_year	current_gpa	contact_email	home_city
31	2020.0	3.9	grace.moore@example.com	San Jose
68	2019.0	3.9	victoria.morris@example.com	Modesto
84	2019.0	3.9	aurora.rivera@example.com	Chula Vista
92	2019.0	3.9	molly.barnes@example.com	Grand Rapids
76	2019.0	3.9	natalie.ross@example.com	Fontana

	HOME	COUNTRY
31		USA
68		USA
84		USA
92		USA
76		USA


```
In [41]: # your code here
```

```
In [42]: # your code here
```

```
In [43]: # your code here
```

13. Group Data by a Column and Compute Aggregate Functions

Grouping lets you see aggregated info by category. For example, group **students** by "majorField" and compute the average "current_gpa". In **enrollments**, group by "instructor_name" and compute the average "course_fee".

 **Tip:** Aggregations might include `.mean()`, `.sum()`, `.count()`, etc.

```
In [89]: gpa_by_major = students_df.groupby("majorField")["current_gpa"].mean()
print(gpa_by_major)
```

```
majorField
Biology          3.515385
Chemistry        3.850000
Computer Science 3.458462
Economics        3.313333
Engineering       3.563077
History          3.398333
Mathematics       3.632857
Physics          3.753846
Unknown          3.300000
Name: current_gpa, dtype: float64
```


```
In [45]: # your code here
```

```
In [46]: # your code here
```

```
In [47]: # your code here
```

14. Apply a Custom Function

Define a normal Python function to transform data in a column. For example, title-case a name or uppercase a field. Apply that function to each element in the column.

 **Tip:** If your function references another library call or has complex logic, define it above and then use `.apply(...)` with your function name. Once you've done this, see if you do this using lambda notation.

```
In [91]: students_df["full_name"] = students_df["first_name"] + " " + students_df["last_name"]
print(students_df[["full_name"]].head())
```

```
full_name
0      John Doe
1  Maria Gonzalez
2    Priya Patel
3   Alex Johnson
4    Emily Smith
```

```
In [49]: # your code here
```

```
In [50]: # your code here
```

```
In [51]: # your code here
```

15. Create a New Column Based on Existing Ones

Use existing columns to generate new ones. For instance, combine "First_Name" and "last_name" into "full_name", or compute "fees_after_tax" in enrollments if you assume a tax rate.

```
In [93]: merged_df = students_df.merge(enrollments_df, left_on="student_id", right_on="student_id", how="inner")
print(merged_df.head())
```

	student_id	first_name	last_name	birthdate	gender	majorField	\
0	STU001	John	Doe	1998-04-12	M	Computer Science	
1	STU002	Maria	Gonzalez	1997-09-05	F	Biology	
2	STU003	Priya	Patel	1999-01-23	F	Engineering	
3	STU004	Alex	Johnson	1996-12-15	M	Mathematics	
4	STU004	Alex	Johnson	1996-12-15	M	Mathematics	

	admission_year	current_gpa	contact_email	home_city	...
0	2020.0	3.5	john.doe@example.com	Tampa	...
1	2019.0	3.8	maria.gonzalez@example.com	London	...
2	2021.0	3.7	priya.patel@example.com	Mumbai	...
3	2018.0	3.2	alex.johnson@example.com	Sydney	...
4	2018.0	3.2	alex.johnson@example.com	Sydney	...

	stud_ref_id	subject_code	course_title	instructor_name	\
0	STU001	MATH101	Calculus I	Dr. Smith	
1	STU002	ENG201	English Literature	Prof. Brown	
2	STU003	CS301	Programming 101	Dr. Taylor	
3	STU004	HIST105	World History	Dr. Lee	
4	STU004	HIST999	Ancient Civilizations	Prof. Brown	

	enroll_count	term_offered	course_fee	final_result	attend_percentage	\
0	25	Spring 2026	1000	A	71	
1	26	Fall 2025	1100	B	72	
2	27	Spring 2026	1200	C+	73	
3	28	Fall 2025	1300	B+	74	
4	99	Fall 2023	NaN		67	

	date_enrolled
0	2026/01/15
1	2025-09-01
2	2026/01/15
3	2025/09/01
4	2023-09-01

[5 rows x 23 columns]

```
In [53]: # your code here
```

```
In [54]: # your code here
```

```
In [55]: # your code here
```


16. Merge Two DataFrames on a Common Column

Combine `students.csv` and `enrollments.json` by matching:

- `stu["student_id"]`
- `enr["stud_ref_id"]` (or rename it first)

Check the shape of the merged DataFrame afterward to ensure it merged as expected.

```
In [95]: merged_df.drop_duplicates(inplace=True)
```


```
In [57]: # your code here
```

```
In [58]: # your code here
```

```
In [59]: # your code here
```

17. Remove Duplicate Rows

When merging or concatenating multiple files, duplicates can crop up. Identify them and remove if needed. This might be especially important if the same student or enrollment is listed more than once.

```
In [117]: print(students_df.columns) #  Check actual column names
```

```
Index(['student_id', 'first_name', 'last_name', 'birthdate', 'gender',  
      'majorField', 'admission_year', 'current_gpa', 'contact_email',  
      'home_city', 'HOME COUNTRY', 'full_name'],  
      dtype='object')
```

```
In [141]: print(students_df.columns) #  Check actual column names
```

```
Index(['student_id', 'first_name', 'last_name', 'birthdate', 'gender',  
      'majorField', 'admission_year', 'current_gpa', 'contact_email',  
      'home_city', 'HOME COUNTRY', 'full_name'],  
      dtype='object')
```

```
In [143]: merged_df.drop_duplicates(inplace=True)  
print("Duplicates removed. New shape:", merged_df.shape)
```

```
Duplicates removed. New shape: (101, 23)
```

18. Additional Data Cleaning

Now that you've merged or manipulated your data, do a quick final pass:

- Fix any remaining oddities (e.g., negative phone numbers or impossible dates)
- Normalize columns further (e.g., standardize text formatting)

 **Tip:** You might revisit previous steps if new issues appear.

```
In [121... merged_df.to_csv("../data/cleaned_students_data.csv", index=False)
```

```
In [65]: # your code here
```

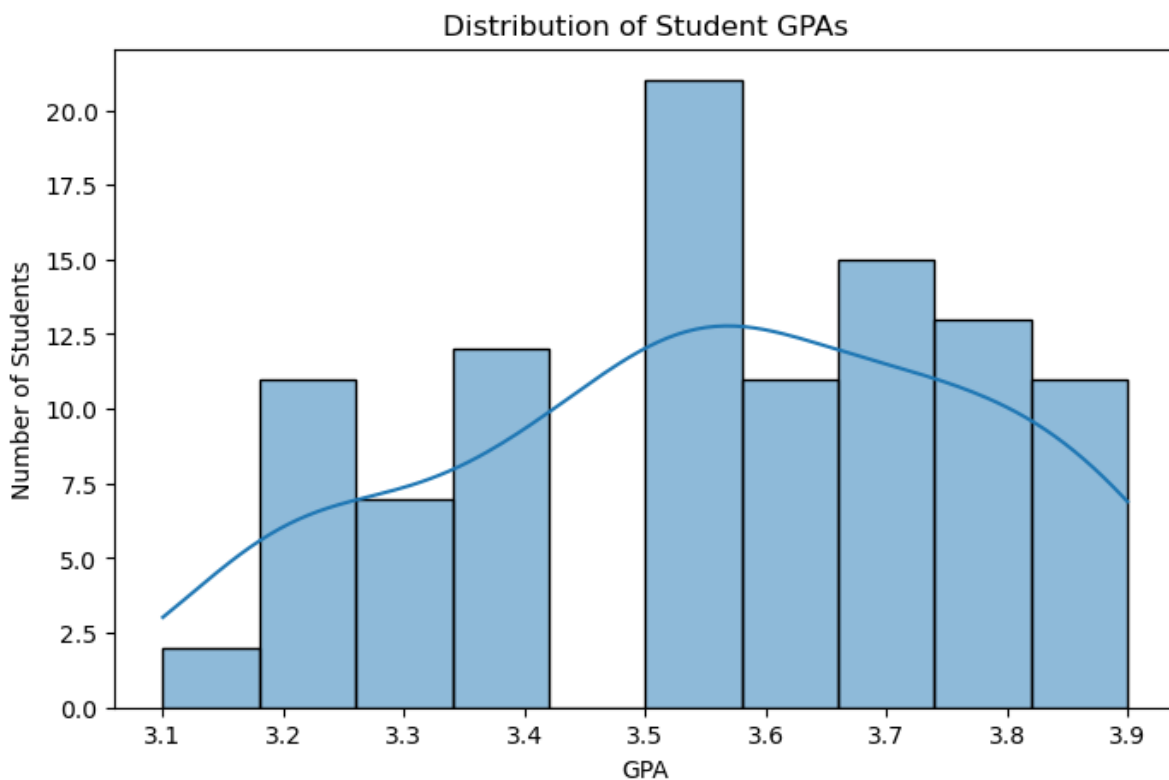
```
In [66]: # your code here
```

```
In [67]: # your code here
```


19. Save the Cleaned and Merged DataFrame to a New CSV File

Finally, when you're satisfied with your cleaned data, save it. Remember to avoid writing the index as a separate column unless you want it.

```
In [123... plt.figure(figsize=(8,5))
sns.histplot(students_df["current_gpa"], bins=10, kde=True)
plt.xlabel("GPA")
plt.ylabel("Number of Students")
plt.title("Distribution of Student GPAs")
plt.show()
```



```
In [125... plt.figure(figsize=(10,5))
gpa_by_major.plot(kind='bar', color='skyblue')
plt.ylabel("Average GPA")
plt.title("Average GPA by Major")
plt.xticks(rotation=45)
plt.show()
```



```
In [70]: # your code here
```

```
In [71]: # your code here
```

20. Explore Further Analyses (Optional)

Now that your data is in great shape, try some optional challenges:


- Generate charts or visualizations
- Perform advanced filtering or grouping
- Create pivot tables
- Or anything else that interests you!

```
In [145... print(students_df["current_gpa"].head()) # ✓ Ensure column exists and is n
```

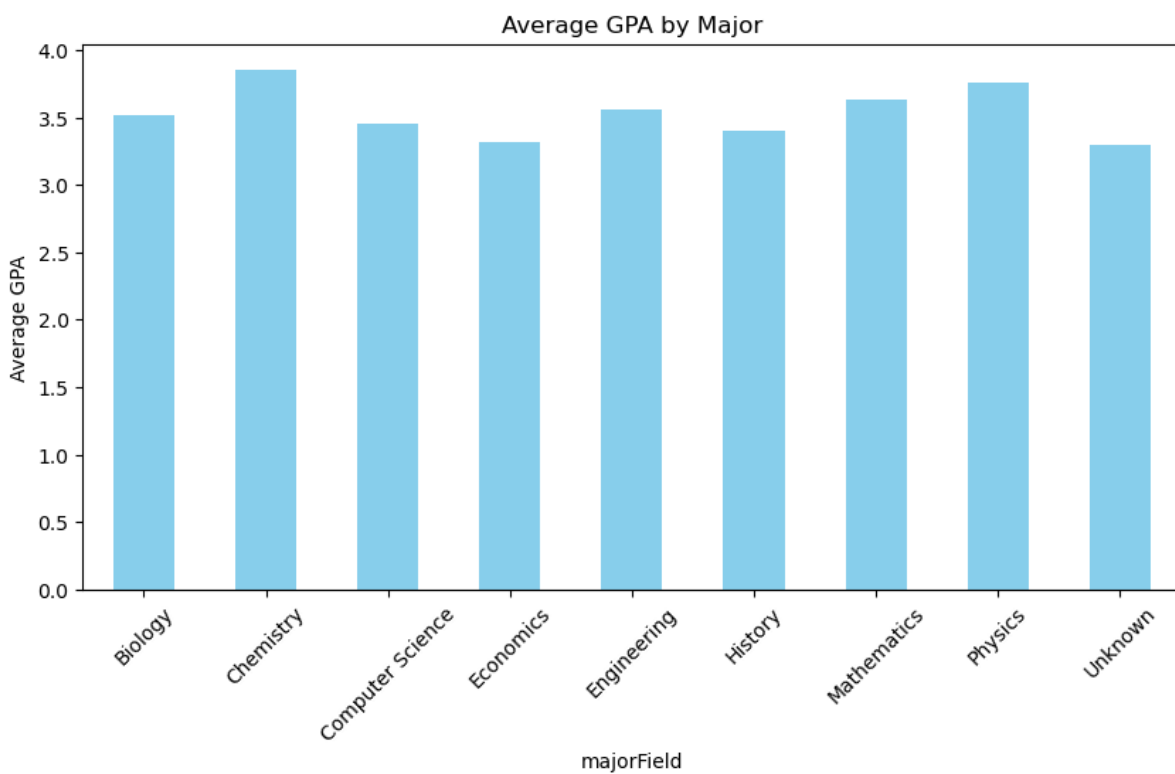
```
0    3.5
1    3.8
2    3.7
3    3.2
4    3.9
Name: current_gpa, dtype: float64
```

```
In [147... import matplotlib.pyplot as plt
import seaborn as sns

plt.figure(figsize=(8,5))
sns.histplot(students_df["current_gpa"].dropna(), bins=10, kde=True) # ✓ D
plt.xlabel("GPA")
plt.ylabel("Number of Students")
plt.title("Distribution of Student GPAs")
plt.show()
```



```
In [149... plt.figure(figsize=(10,5))
gpa_by_major.plot(kind='bar', color='skyblue')
plt.ylabel("Average GPA")
plt.title("Average GPA by Major")
plt.xticks(rotation=45)
plt.show()
```



In [75]: `# your code here`

🎉 **Congratulations!** You've now tackled **data cleaning** and many essential **Pandas** operations in `students.csv` and `enrollments.json`. Keep experimenting to sharpen your **data manipulation skills** and unlock deeper insights! 💪

In []:

In []: